

Parallelising your OCaml code with Multicore OCaml

Sadiq Jaffer, Tom Kelly, Sudha Parimala, KC
Sivaramakrishnan, Anil Madhavapeddy

Overview

- Multicore OCaml
- Domains
- Domainslib
- Further optimisation

Multicore OCaml

- Concurrency is overlapping computations
- Parallelism is simultaneous computations

Multicore OCaml \Rightarrow concurrency *and* shared-memory
parallelism

Multicore OCaml

- Compatible with existing OCaml code (inc ppx)
- OCaml 5 will have parallelism via Domains
- Concurrency via effects and fibers to follow

Domains

- Unit of parallelism
- Heavyweight
- Functionality
 - Spawn/join
 - Wait/notify
 - Atomic memory operations
 - Local storage

N-Body

- Derived from benchmarks game
- Models orbit of a number of bodies



N-Body serial

```
1 let advance bodies n_bodies dt =
2   for i = 0 to n_bodies - 1 do
3     let b = bodies.(i) in
4     for j = i+1 to n_bodies - 1 do
5       let b_o = bodies.(j) in
6       let dx = b.x -. b_o.x and dy = b.y -. b_o.y
7         and dz = b.z -. b_o.z in
8       let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
9       let mag = dt /. (dist2 *. sqrt(dist2)) in
10
11       b.vx <- b.vx -. dx *. b_o.mass *. mag;
12       b.vy <- b.vy -. dy *. b_o.mass *. mag;
13       b.vz <- b.vz -. dz *. b_o.mass *. mag;
14
15       b_o.vx <- b_o.vx +. dx *. b.mass *. mag;
16       b_o.vy <- b_o.vy +. dy *. b.mass *. mag;
```

N-Body serial

```
1 let advance bodies n_bodies dt =
2   for i = 0 to n_bodies - 1 do
3     let b = bodies.(i) in
4     for j = i+1 to n_bodies - 1 do
5       let b_o = bodies.(j) in
6       let dx = b.x -. b_o.x and dy = b.y -. b_o.y
7         and dz = b.z -. b_o.z in
8       let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
9       let mag = dt /. (dist2 *. sqrt(dist2)) in
10
11       b.vx <- b.vx -. dx *. b_o.mass *. mag;
12       b.vy <- b.vy -. dy *. b_o.mass *. mag;
13       b.vz <- b.vz -. dz *. b_o.mass *. mag;
14
15       b_o.vx <- b_o.vx +. dx *. b.mass *. mag;
16       b_o.vy <- b_o.vy +. dy *. b.mass *. mag;
```


N-Body serial

```
4   for j = 1+1 to n_bodies - 1 do
5       let b_o = bodies.(j) in
6       let dx = b.x -. b_o.x and dy = b.y -. b_o.y
7           and dz = b.z -. b_o.z in
8       let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
9       let mag = dt /. (dist2 *. sqrt(dist2)) in
10
11       b.vx <- b.vx -. dx *. b_o.mass *. mag;
12       b.vy <- b.vy -. dy *. b_o.mass *. mag;
13       b.vz <- b.vz -. dz *. b_o.mass *. mag;
14
15       b_o.vx <- b_o.vx +. dx *. b.mass *. mag;
16       b_o.vy <- b_o.vy +. dy *. b.mass *. mag;
17       b_o.vz <- b_o.vz +. dz *. b.mass *. mag;
18   done
19 done
```

N-Body serial

All experiments on an 2x Xeon E5-2695 v4

```
real    1m23.423s  
user    1m23.422s  
sys     0m0.000s
```

(256 iterations, 8192 bodies)

How fast can we go?

Amdahl's law for parallel programs:

$$\frac{1}{(1 - p) + \left(\frac{p}{s}\right)}$$

p = proportion of parallelisable code

s = degree of parallelism

Linux Perf

Sampling profile of `nbody_serial.exe` run:

Children	Self	Command	Shared Object	Symbol
+ 99.95%	0.00%	nbody_serial.exe	nbody_serial.exe	[.] caml_start_program
+ 99.95%	0.00%	nbody_serial.exe	nbody_serial.exe	[.] caml_program
+ 99.95%	0.00%	nbody_serial.exe	nbody_serial.exe	[.] camlDune__exe__Nbody_serial__entry
+ 99.55%	99.55%	nbody_serial.exe	nbody_serial.exe	[.] camlDune__exe__Nbody_serial__advance_90
0.34%	0.34%	nbody_serial.exe	nbody_serial.exe	[.] camlDune__exe__Nbody_serial__energy_159
0.05%	0.05%	nbody_serial.exe	nbody_serial.exe	[.] camlDune__exe__Nbody_serial__update_152
0.03%	0.00%	nbody_serial.exe	nbody_serial.exe	[.] _start

Perfect scalability \Rightarrow $\sim 220x$ speedup

If $p < 0.89$, max speedup single digits!

Iteration 1: Domain per body

Don't do this

```
1 let advance bodies n_bodies dt =
2   let ds =
3     Array.init n_bodies (fun i -> Domain.spawn (fun _ ->
4       let b = bodies.(i) in
5       for j = 0 to n_bodies - 1 do
6         let b_o = bodies.(j) in
7         if (i!=j) then begin
8           let dx = b.x -. b_o.x and dy = b.y -. b_o.y
9           and dz = b.z -. b_o.z in
10          let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
11          let mag = dt /. (dist2 *. sqrt(dist2)) in
12          b.vx <- b.vx -. dx *. b_o.mass *. mag;
13          b.vy <- b.vy -. dy *. b_o.mass *. mag;
14          b.vz <- b.vz -. dz *. b_o.mass *. mag;
15        end
16      done
```

Iteration 1: Domain per body

Don't do this

```
2  let ds =
3    Array.init n_bodies (fun i -> Domain.spawn (fun _ ->
4      let b = bodies.(i) in
5      for j = 0 to n_bodies - 1 do
6        let b_o = bodies.(j) in
7        if (i!=j) then begin
8          let dx = b.x -. b_o.x and dy = b.y -. b_o.y
9          and dz = b.z -. b_o.z in
10         let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
11         let mag = dt /. (dist2 *. sqrt(dist2)) in
12         b.vx <- b.vx -. dx *. b_o.mass *. mag;
13         b.vy <- b.vy -. dy *. b_o.mass *. mag;
14         b.vz <- b.vz -. dz *. b_o.mass *. mag;
15       end
16     done
17   )) in
```

Iteration 1: Domain per body

Don't do this

```
3   Array.init n_bodies (fun i -> Domain.spawn (fun _ ->
4       let b = bodies.(i) in
5       for j = 0 to n_bodies - 1 do
6           let b_o = bodies.(j) in
7           if (i!=j) then begin
8               let dx = b.x -. b_o.x and dy = b.y -. b_o.y
9               and dz = b.z -. b_o.z in
10              let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
11              let mag = dt /. (dist2 *. sqrt(dist2)) in
12              b.vx <- b.vx -. dx *. b_o.mass *. mag;
13              b.vy <- b.vy -. dy *. b_o.mass *. mag;
14              b.vz <- b.vz -. dz *. b_o.mass *. mag;
15          end
16      done
17  )) in
18  Array.iter (Domain.join) ds
```

Iteration 1: Domain per body

Oops.

```
real    8m10.965s
user    25m24.372s
sys     11m30.816s
```

Domains are heavyweight
Aim for same number as cores
Spawn/join infrequently

Domainslib

- Task pool
 - Parallel
 - for / reduce / scan
 - Async/await
- Channels

<https://github.com/ocaml-multicore/domainslib/>

Iteration 2: Domainslib

```
1 let advance pool n_bodies n_domains bodies dt =  
2   T.parallel_for pool  
3     ~chunk_size:(n_bodies/n_domains)  
4     ~start:0  
5     ~finish:(n_bodies - 1)  
6     ~body:(fun i ->  
7       let b = bodies.(i) in  
8       for j = 0 to n_bodies - 1 do  
9         let b_o = bodies.(j) in  
10        if (i!=j) then begin  
11          let dx = b.x -. b_o.x and dy = b.y -. b_o.y  
12          and dz = b.z -. b_o.z in  
13          let dist2 = dx *. dx +. dy *. dy +. dz *. dz in  
14          let mag = dt /. (dist2 *. sqrt(dist2)) in  
15          b.vx <- b.vx -. dx *. b_o.mass *. mag;  
16          b.vy <- b.vy -. dy *. b_o.mass *. mag;
```

Iteration 2: Domainslib

```
1  scale: 1
5  ~finish:(n_bodies - 1)
6  ~body:(fun i ->
7    let b = bodies.(i) in
8    for j = 0 to n_bodies - 1 do
9      let b_o = bodies.(j) in
10     if (i!=j) then begin
11       let dx = b.x -. b_o.x and dy = b.y -. b_o.y
12       and dz = b.z -. b_o.z in
13       let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
14       let mag = dt /. (dist2 *. sqrt(dist2)) in
15       b.vx <- b.vx -. dx *. b_o.mass *. mag;
16       b.vy <- b.vy -. dy *. b_o.mass *. mag;
17       b.vz <- b.vz -. dz *. b_o.mass *. mag;
18     end
19   done
```

Iteration 2: Domainslib

Cores	Time	Vs Serial	Vs Self
1	128.076s	0.65x	1x
2	54.987s	1.51x	2.32x
3	42.577s	1.94x	3.00x
4	32.753s	2.53x	3.90x
8	18.868s	4.39x	6.78x
16	11.438s	7.28x	11.19x
24	8.465s	9.80x	15.12x

Shared state pitfalls

Minimise writes to frequently read shared state

Avoid contended writes to shared state

Use `perf stat`, `perf record` and `perf c2c`

```
perf stat -e cycles,cycle_activity.cycles_no_execute,cycle_acti  
  
458,432,211,938      cycles  
74,023,598,176      cycle_activity.cycles_no_execute  
67,741,246,568      cycle_activity.stalls_mem_any  
481,904,795          cache-misses  
14,350,052           mem_load_uops_l3_miss_retired.remote_hitm  
  
8.726805280 seconds time elapsed
```

Iteration 3: Avoiding contention

```
1 type planet_pos = { mutable x : float; mutable y : float; mu
2 type planet_vec = { mutable vx: float; mutable vy: float; mu
3
4 let advance pool n_domains n_bodies bodies_pos bodies_vec dt =
5   T.parallel_for pool
6     ~chunk_size:(n_bodies/n_domains)
7     ~start:0
8     ~finish:(n_bodies - 1)
9     ~body:(fun i ->
10       let bp = bodies_pos.(i) in
11       let bv = bodies_vec.(i) in
12       let vx, vy, vz = ref bv.vx, ref bv.vy, ref bv.vz in
13       for j = 0 to n_bodies - 1 do
14         let bp' = bodies_pos.(j) in
15         if (i!=j) then begin
16           let dx = bp.x - bp'.x and dy = bp.y - bp'.y and
```

Iteration 3: Avoiding contention

```
1  ~advance_pool: n_domains n_bodies bodies_pos bodies_vec do
5    T.parallel_for pool
6      ~chunk_size:(n_bodies/n_domains)
7      ~start:0
8      ~finish:(n_bodies - 1)
9      ~body:(fun i ->
10         let bp = bodies_pos.(i) in
11         let bv = bodies_vec.(i) in
12         let vx, vy, vz = ref bv.vx, ref bv.vy, ref bv.vz in
13         for j = 0 to n_bodies - 1 do
14           let bp' = bodies_pos.(j) in
15           if (i!=j) then begin
16             let dx = bp.x -. bp'.x and dy = bp.y -. bp'.y and
17             let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
18
19             let mag = dt /. (dist2 *. sqrt(dist2)) in
20             let mass = bp'.mass in
```


Iteration 3: Avoiding contention


```
12 let vx, vy, vz = let bv.vx, let bv.vy, let bv.vz in
13 for j = 0 to n_bodies - 1 do
14   let bp' = bodies_pos.(j) in
15   if (i!=j) then begin
16     let dx = bp.x -. bp'.x and dy = bp.y -. bp'.y and
17     let dist2 = dx *. dx +. dy *. dy +. dz *. dz in
18
19     let mag = dt /. (dist2 *. sqrt(dist2)) in
20     let mass = bp'.mass in
21     vx := !vx -. dx *. mass *. mag;
22     vy := !vy -. dy *. mass *. mag;
23     vz := !vz -. dz *. mass *. mag;
24   end
25 done;
26 bv.vx <- !vx;
27 bv.vy <- !vy;
```


Iteration 3: Avoiding contention

```
perf stat -e cycles,cycle_activity.cycles_no_execute,cycle_activi  
  
341,067,775,833      cycles  
31,523,253,083      cycle_activity.cycles_no_execute  
29,569,589,458      cycle_activity.stalls_mem_any  
7,009,375           cache-misses  
956,792             mem_load_uops_l3_miss_retired.remote_hitm  
  
6.639116502 seconds time elapsed
```

~35% speedup at 24 cores





Take aways

- Multicore is ready to use, upstreaming in progress
- Profile to understand serial performance
- Use Domainslib abstractions to add parallelism
- Share work as coarsely as you can
- Avoid writing to shared state as much as possible

Use the [multicore github issue tracker](#) if you find unexpected behaviour